REPUBLIC OF TURKEY CUKUROVA UNIVERSITY FACULTY OF ENGINEERING
DEPARTMENT OF COMPUTER ENGINEERING

BANDWIDTH OPTIMIZATION IN VIDEO TRANSMISSION VIA DEEP LEARNING-BASED 'SPLIT COMPUTING'

GRADUATION PROJECT FINAL REPORT

Prepared by: Halit ŞEN 2021555060

Advisor: Assoc. Prof. Dr. Mehmet SARIGÜL

ADANA – 2026

CONTENTS

ABSTRACT..... i

1. INTRODUCTION AND MOTIVATION 1
   1.1. Objective of the Project and the Bandwidth Problem 1
   1.2. Split Computing and Asymmetric Network Philosophy 2
   1.3. Organization of the Report 3
2. SPATIAL ENHANCEMENT: SUPER RESOLUTION (ESPCN) 4
   2.1. Feature Extraction in Low-Resolution Space 4
   2.2. Luma (Y) Channel Optimization and Hybrid Approach 5
   2.3. Hardware Performance (Apple M1) and Real-Time Synthesis 7
3. TEMPORAL ENHANCEMENT: VIDEO FRAME INTERPOLATION (VFI) 9
   3.1. LiteVFINet and Hierarchical Multi-Frame Prediction 9
   3.2. Patch-Based Architecture: MicroVFINet 10
   3.3. CoordConv (Spatial Awareness) and Sobel Edge Loss 12
   3.4. Optical Flow-Aware Autonomous Data Filtering 14
4. ASYMMETRIC BOTTLENECK ARCHITECTURE AND COMPRESSION 17
   4.1. Evaluation Criterion: Standard ZLIB vs. AI-Assisted ZLIB 17
   4.2. Bottleneck Quantization and Dictionary Compatibility (V3 - V5) 18
   4.3. Autonomous Self-Checking and Bandwidth Savings 20
   4.4. Visual Analyses and the Color Loss Discussion 22
5. EXPLORING THE LIMITS: ADVANCED BOTTLENECK AND ENGINEERING LESSONS 24
   5.1. Color Fading Crisis and the "Biological Isolation" Lesson (YUV) 24
   5.2. Capacity Paradox and Model Stubbornness (Reward Hacking) 26
   5.3. Dimensional Mismatch and Volume Conservation 28
   5.4. The Ultimate Crisis: "Dimension Paradox" and Shaping the Future Vision .. 30

ABSTRACT

In this study, an end-to-end hybrid deep learning system based on the "Split Computing" paradigm has been developed to overcome the bandwidth constraints encountered in high-resolution video transmission. In scenarios where traditional compression algorithms alone are insufficient, spatial and temporal enhancements were performed on videos using the local hardware resources of the client device (such as the Apple M1 GPU) to alleviate the data transmission load.

The system consists of two fundamental hardware-friendly modules: the hybrid ESPCN architecture, which focuses solely on the luminance (Luma) channel to increase spatial resolution; and the patch-based, optical flow-aware MicroVFINetPro architecture, designed to reduce high processing loads while improving temporal fluency (FPS). To physically reduce the data transmission cost between the server and the client, an asymmetric U-Net structure was utilized, and all data transmission was intended to be compressed into a highly narrow "Bottleneck" layer. The tensor data passing through the bottleneck were quantized into an 8-bit format using STE algorithms to ensure full compatibility with lossless compression algorithms like ZLIB. An "Autonomous Self-Checking" mechanism was developed to prevent the network from transmitting certain regions unnecessarily, enabling the model to determine its own penalty coefficient based on image complexity.

Experimental findings demonstrate that the developed architecture provides a net data saving of approximately 20% in network traffic without applying any spatial reduction (stride) operation. The client-side models can successfully synthesize at high speeds and with high structural accuracy (28-29 dB PSNR) on Apple M1 hardware. This developed system reveals that artificial intelligence can be utilized as an autonomous smart filter that maximizes the efficiency of traditional methods in data transmission, rather than completely replacing them.

Keywords: Video Transmission, Bandwidth Optimization, Split Computing, Super Resolution, ESPCN, Video Frame Interpolation, Quantization, Bottleneck.

1. INTRODUCTION AND MOTIVATION

1.1. Objective of the Project and the Bandwidth Problem

Nowadays, the rapid proliferation of digital broadcasting, cloud gaming, and video conferencing systems has increased the data traffic load on network infrastructures to unprecedented levels. In particular, the rising demand for ultra-high-definition and high frame rate (FPS) video content, such as 4K/8K, is pushing the limits of existing internet bandwidth. While traditional video compression standards attempt to manage this

heavy data flow, they often lead to visual quality loss at high compression ratios or undesirable delays due to increased computational costs.

The primary objective of this study is to minimize the bandwidth constraints encountered during the transmission of high-quality video data over the network using deep learning methods. To this end, an end-to-end hybrid system has been developed that enables the real-time enhancement of videos—which are intentionally transmitted by the server at low resolution and limited frame rates—using artificial intelligence algorithms on the client's local hardware resources. Thus, the goal is to provide users with a seamless and high-quality viewing experience while relieving network traffic of unnecessary loads.

1.2. Split Computing and Asymmetric Network Philosophy

The architectural foundation of the project is built upon the "Split Computing" paradigm, which aims to transcend the boundaries of traditional purely cloud-based or purely edge-device-based processing. The split computing approach aims to logically distribute the computational load between the server and the client ecosystem.

In this context, an original asymmetric network design (Encoder-Decoder) has been adopted, inspired by the standard U-Net architecture in the literature, but abandoning its symmetric structure and intensive information transfers (skip connections) that do not align with the low-bandwidth vision of the project. Heavy feature extraction and computational load are assigned to the encoder block running on the server side. The decoder network, which will operate on the client device (Apple M1, etc.), is designed to be extremely shallow and lightweight to perform hardware-friendly and real-time inference.

To minimize the transmission cost over the network, high-cost intermediate information transfers have been cancelled, and all communication of the network has been forced to pass through a single "Bottleneck" layer. Thanks to this asymmetric constraint, the model has learned to physically compress the most important structural details into the smallest possible size and "reconstruct the highest visual quality while consuming the least bytes on the client side," rather than memorizing the data.

1.3. Organization of the Report (Raporun Organizasyonu)

This final report is structured to present the theoretical infrastructure, development phases (ablation studies), and final experimental findings of the developed hybrid video transmission system within a systematic academic framework. The sections of the report are organized as follows:

- **Section 2** analyzes the spatial super-resolution (ESPCN) models developed to enhance low-resolution images on the client side in a hardware-friendly manner and the resulting hardware acceleration.

- **Section 3** examines the patch-based MicroVFINetPro architecture designed to increase the temporal fluidity of videos, optical flow-based data filtering processes, and multi-frame prediction algorithms.
- **Section 4** addresses the asymmetric bottleneck structures that provide physical communication between the server and the client with the Split Computing approach, ZLIB-compatible STE-based 8-bit quantization operations, and the Autonomous Self-Checking mechanism developed against the "Reward Hacking" crisis.
- **Section 5** presents the Luma-Chroma (SoC) separation, the grid-based autonomous fusion system, and the applications of transitioning from patch space to video space, which constitute the most recent research phase of the project.
- **Section 6** presents the test results of all developed architectures on local hardware, metric analyses (PSNR/SSIM, FPS), and bandwidth savings , thereby determining the mathematical boundaries of the dynamic server-client division of labor.
- **Section 7** provides a general evaluation in light of the scientific inferences obtained and discusses the vision offered by the project for future communication systems.

2. SPATIAL ENHANCEMENT: SUPER RESOLUTION (ESPCN)

Enhancing the spatial quality of videos is one of the system's first and most critical tasks on the client side. At this stage, the fundamental problem is the real-time conversion of missing high-frequency details (edges, textures) in low-resolution (LR) video frames into high-resolution (HR) format on limited hardware resources. Accordingly, the ESPCN (Efficient Sub-Pixel Convolutional Neural Network) architecture, which stands out due to its hardware-friendly structure, was taken as a basis and optimized according to system requirements.

2.1. Feature Extraction in Low-Resolution Space

Traditional super-resolution networks (e.g., SRCNN) scale the image to the target size using conventional interpolation methods (Bicubic, etc.) before feeding it into the network and perform all convolution operations in this high-resolution space. Since this approach exponentially increases the computational cost (FLOPs) and memory consumption, it is not suitable for the goal of real-time video processing.

In the developed ESPCN-based architecture, feature extraction layers were run directly in the low-resolution (LR) space. The convolution operations forming the core of the

network were performed on smaller-sized tensors, thereby minimizing the hardware load. Learned deep spatial relationships are converted directly into HR format through a mathematical reshaping operation, thanks to the "Sub-Pixel Convolution" (Pixel Shuffle) layer located at the final stage of the network. This architectural choice allowed blurry or compressed data to reach a clear and sharp form without the need for processor-intensive operations such as deconvolution.

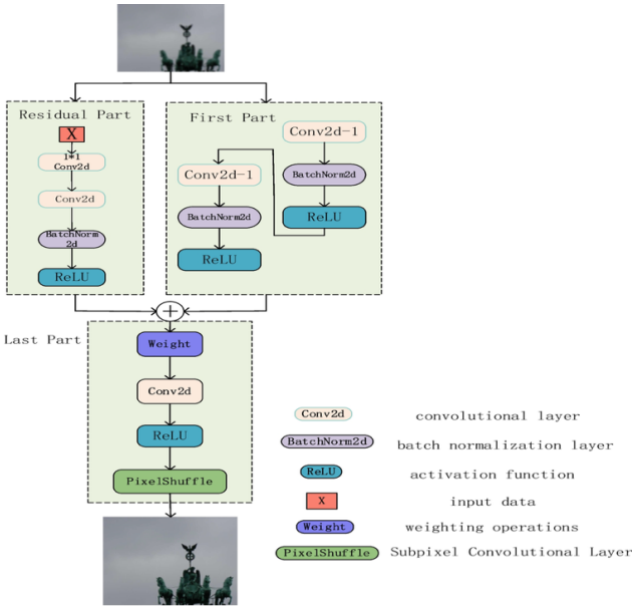

Figure 2.1: Feature extraction in low-resolution space and the Pixel Shuffle layer in the developed ESPCN architecture.

2.2. Luma (Y) Channel Optimization and Hybrid Approach

In the initial ablation studies conducted to increase network capacity (Model 2), increasing the number of filters to 128 and adding attention mechanisms increased the processing time to ~254 milliseconds per frame, dropping the system far behind real-time limits. At this point, based on the principle that more parameters do not always yield better results, a "Hybrid Video Enhancement" strategy was developed that dramatically lightens the hardware load.

This is based on the biological fact that the human eye shows high sensitivity to structural details and light intensity (Luminance) while being more tolerant of color changes (Chrominance). Input images were converted from the traditional 3-channel RGB format to the YCbCr color space, and only the single-channel Y (Luma) matrix, which contains high-frequency fine details, was fed into the deep learning model. The Cb and Cr color channels, which create computational costs, were kept outside the network, scaled using the traditional Bicubic interpolation method, and recombined at the network output.

2.3. Hardware Performance (Apple M1) and Real-Time Inference

The performance of the developed models on the client side was tested on hardware with Apple M1 architecture using the Metal Performance Shaders (MPS) API. The unified memory architecture of the M1 chip provided an advantage in processing high-resolution images by minimizing data transfer delays (overhead) between the CPU and GPU.

The Hybrid ESPCN model (Model 4), utilizing Y-channel optimization, represents the peak of the hardware acceleration achieved. While the processing time per frame was an average of 52.5 milliseconds in the 3-channel (RGB) structure with an attention mechanism, this time was reduced to the 11.74 - 13.27 millisecond range in the hybrid structure. In this way, low-quality inputs at 640x360 (360p) resolution were scaled to 1920x1080 (1080p) format at an average speed of ~80 FPS (75.38 - 85.16 FPS).

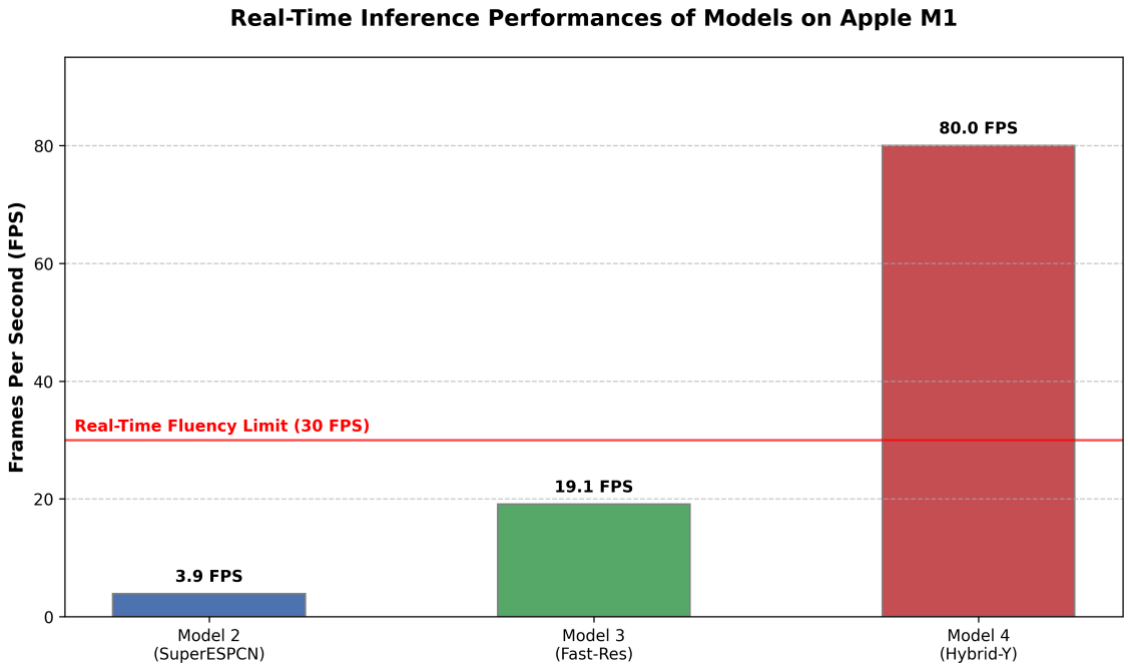

Figure 2.2: Comparison of inference speeds (FPS) of the models (Baseline, Model 2, Fast-Res, and Hybrid) on the Apple M1 architecture.

Despite this massive jump in speed, no critical loss in visual quality occurred. While the model reached 26.66 dB PSNR and 0.817 SSIM values in the evaluation conducted over the Y channel , it demonstrated an average performance of 26.42 dB PSNR in the RGB-combined (Hybrid) results. These results prove that the "divide and conquer" strategy (focusing on the Luma channel) in asymmetric systems is the most optimal method to maximize the performance/cost ratio on edge devices with limited resources.

As visual evidence of the rate-distortion balance, the reconstruction performances of the basic ESPCN architecture and the hardware-accelerated hybrid architecture were compared.

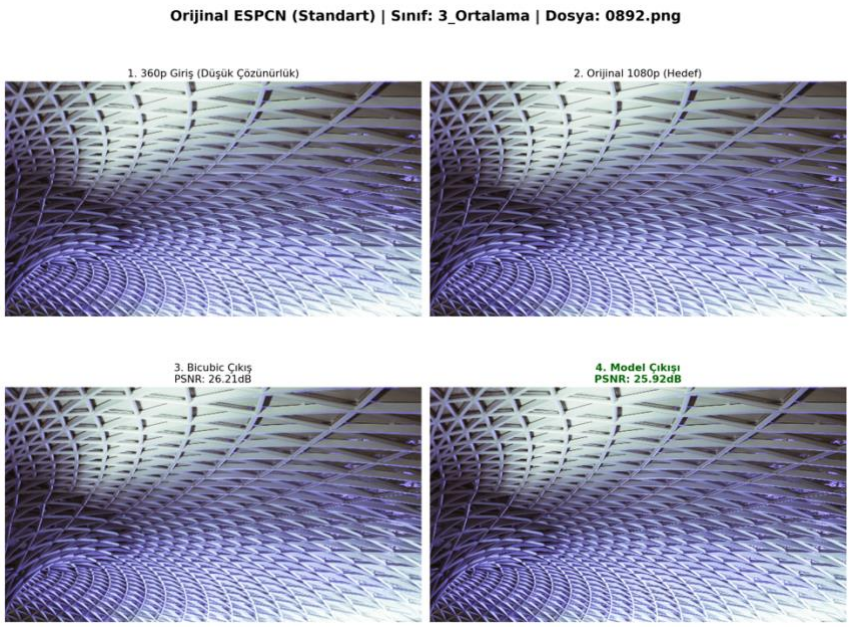

Figure 2.3: Spatial enhancement output obtained using the basic ESPCN network (Average performance example).

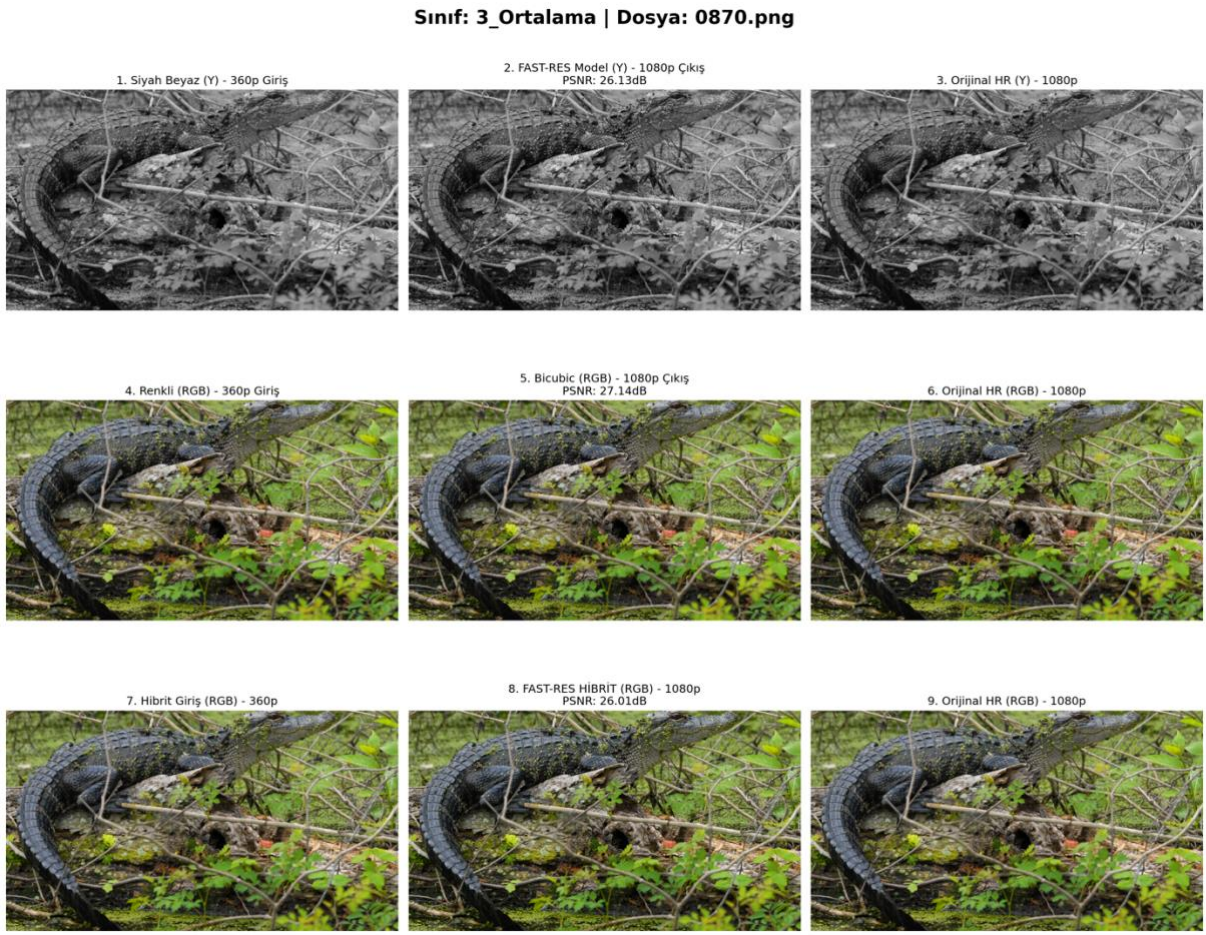

Figure 2.4: Spatial enhancement output obtained with the Y-Channel optimized hybrid architecture (Average performance example).

The visual analyses presented in Figure 2.3 and Figure 2.4 confirm that the radical simplification in the architecture and the single-channel (Luma) strategy do not create a perceptual weakness in terms of edge sharpness and texture integrity compared to the reference model.

3. TEMPORAL ENHANCEMENT: VIDEO FRAME INTERPOLATION (VFI)

Following high-resolution spatial enhancement (Super Resolution), the second critical factor determining user experience in video transmission is temporal fluency (FPS). To preserve bandwidth, a Video Frame Interpolation (VFI) module was developed to smoothly render video streams on the client side, which are intentionally transmitted from the server at low frame rates (e.g., 15 or 30 FPS). The VFI process is not a simple pixel blending; rather, it relies on the physical prediction of objects' positions in intermediate time frames by analyzing optical flow dynamics via artificial intelligence.

3.1. LiteVFINet and Hierarchical Multi-Frame Prediction

In the first stage of the temporal enhancement phase, a U-Net-based LiteVFINet architecture was designed, which is hardware-friendly and capable of end-to-end direct synthesis. To avoid the high computational costs of traditional optical flow networks, the model was trained to take the 1st and 3rd frames directly as inputs and synthesize the missing intermediate frame.

One of the most innovative contributions of this module was the "Modular Hierarchical Scheduling Engine". This engine, capable of synthesizing up to 3 or 7 artificial frames between two consecutive real frames in scenes with intensive linear motion rather than at a fixed ratio, utilizes a "Zero-Latency Cache" strategy. By keeping the centrally calculated frames—which are mandatorily computed in the background during frame generation—in memory, the subsequent frames are projected onto the screen with a 0.0 ms delay, thereby utilizing hardware limits in the most efficient manner.

3.2. Patch-Based Architecture: MicroVFINet

Although LiteVFINet was successful in full-frame predictions, speed bottlenecks began to occur on client hardware when the resolution reached 1080p levels. To overcome this limitation, a radical change was made in the system architecture, abandoning the full-screen processing logic and developing a patch-based "MicroVFINet" architecture that analyzes only the areas where motion is concentrated.

MicroVFINet is an ultra-lightweight network where the number of parameters and channels is dramatically reduced (maximum 64 filters) and convolution-based (stride=2) downsampling is used instead of max-pooling.

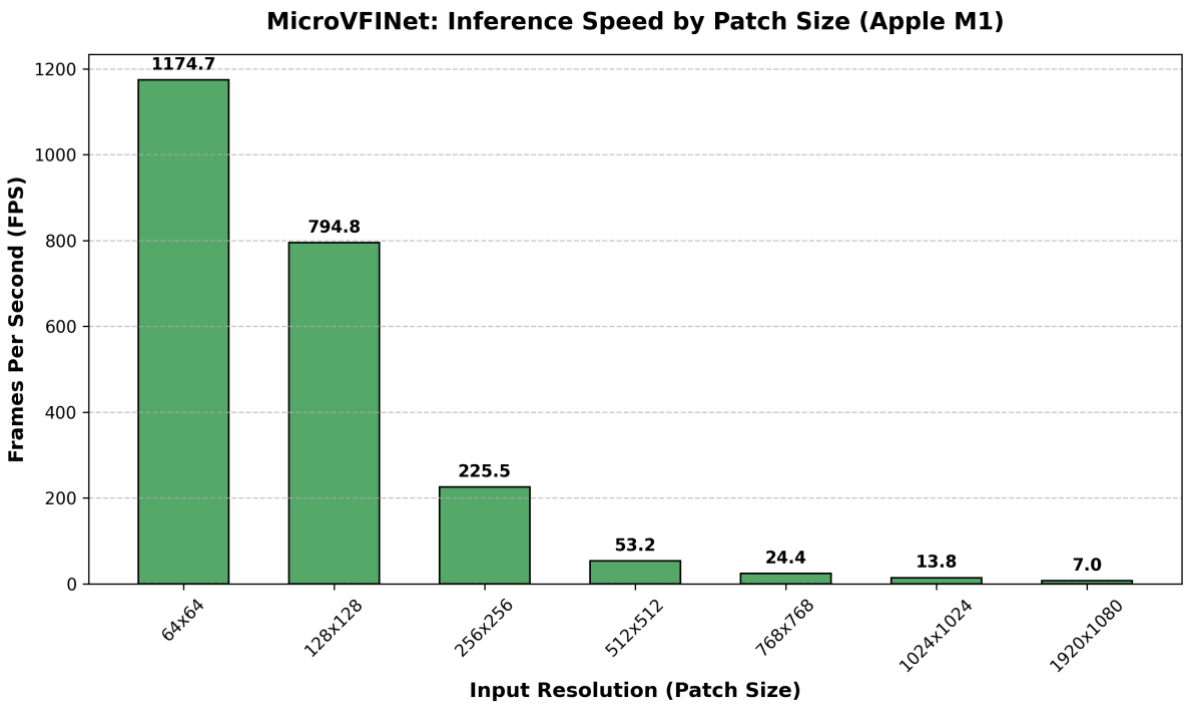

Figure 3.1: Frames per second (FPS) performance of the MicroVFINet architecture on the Apple M1 relative to input resolution (patch size).

3.3. CoordConv (Spatial Awareness) and Sobel Edge Loss

Despite the speed advantage brought by patch-based processing, the translation invariance nature of traditional Convolutional Neural Networks (CNNs) caused a weakness in visual accuracy since the network could not determine where the processed patch originated from on the screen (spatial context).

To solve this problem, the network architecture was updated to MicroVFINetPro, and a "Spatial Awareness" (CoordConv) module was integrated into the system. By adding X and Y coordinate maps to the model as the 7th and 8th channels, the network gained the ability to learn the absolute positions of the pixels. Additionally, to break the over-smoothing effect created by the standard L1 (MAE) loss function on moving edges, "Sobel Edge Loss" was integrated into the system using Sobel matrices that calculate the spatial derivatives of the pixels.

3.4. Optical Flow-Aware Autonomous Data Filtering

Based on the fact that the success of deep learning algorithms depends directly on the quality of the training data, an Optical Flow-based asymmetric data pipeline was constructed instead of feeding random patches to the network.

To push the model's learning capacity, the training set was composed of patches containing "Hard" motions (displacements between 4.0 px and 30.0 px), while the test set was compiled from a broad spectrum containing real-world camera shakes and

micro-motions. The training process was made autonomous using a next-generation "Oscillation Tracker" that mathematically detects when the model reaches saturation (when the loss value gets stuck in a specific band), completely preventing overfitting.

The finalized Model 11 remained in the 26.27 dB PSNR band for macro motions considered "hard" in the literature, while achieving near-perfect frame interpolation by rising to the 29.91 dB PSNR level in scenes containing micro and medium-level optical flow between 0-8 pixels.

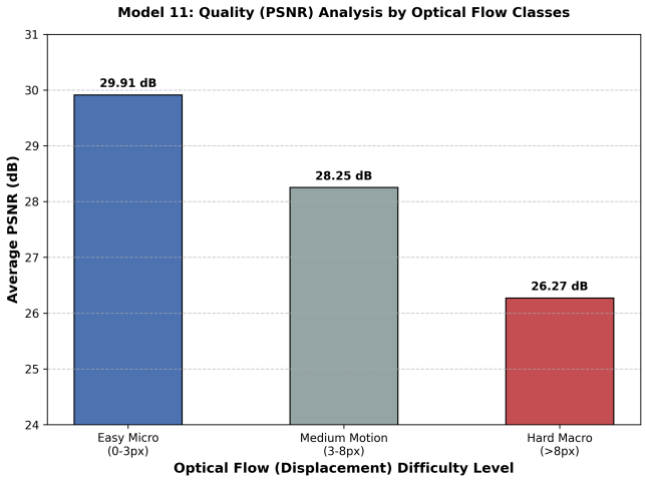

Figure 3.2: PSNR averages achieved by Model 11 according to inter-pixel displacement (Optical Flow) difficulty classes.

The statistical data presented in Figure 3.2 draws a "performance boundary" proving at which motion limits the system should trust the local hardware and in which complex scenes it should download full-resolution data from the server. This statistical finding is also confirmed by visual results:

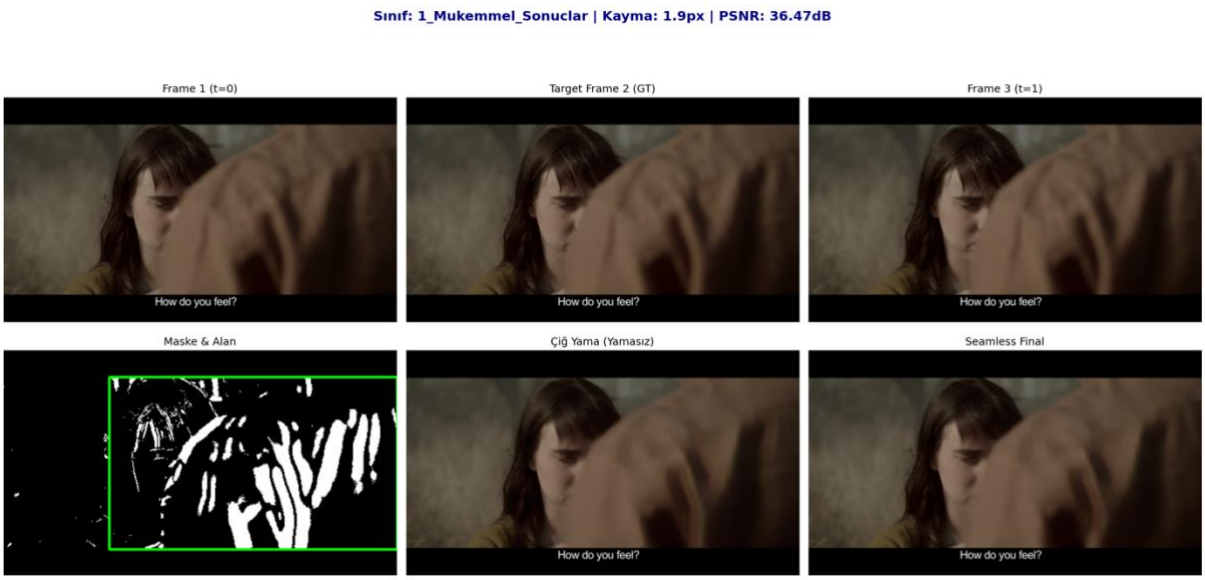

Figure 3.3: High structural accuracy exhibited by the model in a scene containing micro and medium-level optical flow (PSNR: 36.5 dB).

Figure 3.4: Visual distortions experienced by the model in macro motions above 8 pixels (occlusion and severe deformation) (PSNR: 24.2 dB).

As clearly seen in Figure 3.3 and Figure 3.4, while the system successfully preserves edge boundaries and object integrity when staying below the 8-pixel displacement limit, it begins to produce "ghosting" and tearing errors in scenes where objects cross one another or change form once this limit is exceeded. This situation has precisely determined the scientific boundaries of the system's autonomous decision mechanism.

4. ASYMMETRIC BOTTLENECK ARCHITECTURE AND COMPRESSION

The heart of the developed hybrid video transmission system is the "Asymmetric Bottleneck" architecture, which autonomously manages physical communication between the server and the client. While data is enhanced after reaching the client in traditional approaches , in this new phase designed with the "Split Computing" paradigm, the data is physically compressed by neural networks on the server before going out to the network.

4.1. U-Net Inspired Asymmetric Communication Network and Design

The network architecture was inspired by the U-Net encoder-decoder philosophy in the literature. However, the completely symmetric structure of the classic U-Net contradicts the "low bandwidth and hardware-friendly" vision of the project.

Since there is no luxury of sending extra data to the client over the network, high-cost information transfers (skip connections) were completely cancelled, and the network was asymmetrically redesigned. Heavy feature extraction is assigned to the "Bulk Encoder" block running on the server side , while the decoder network, which will operate on edge devices like the Apple M1 , is kept extremely shallow (parameters-light)

and low computational cost (low-FLOPs) to perform real-time inference. Thanks to this asymmetric constraint, the model learned not to memorize the data, but to physically compress the most important structural details into the smallest possible dimension and to reconstruct the highest visual quality by consuming the least bytes on the client side.

Figure 4.1: End-to-End Asymmetric Bottleneck and Autonomous Masking Architecture

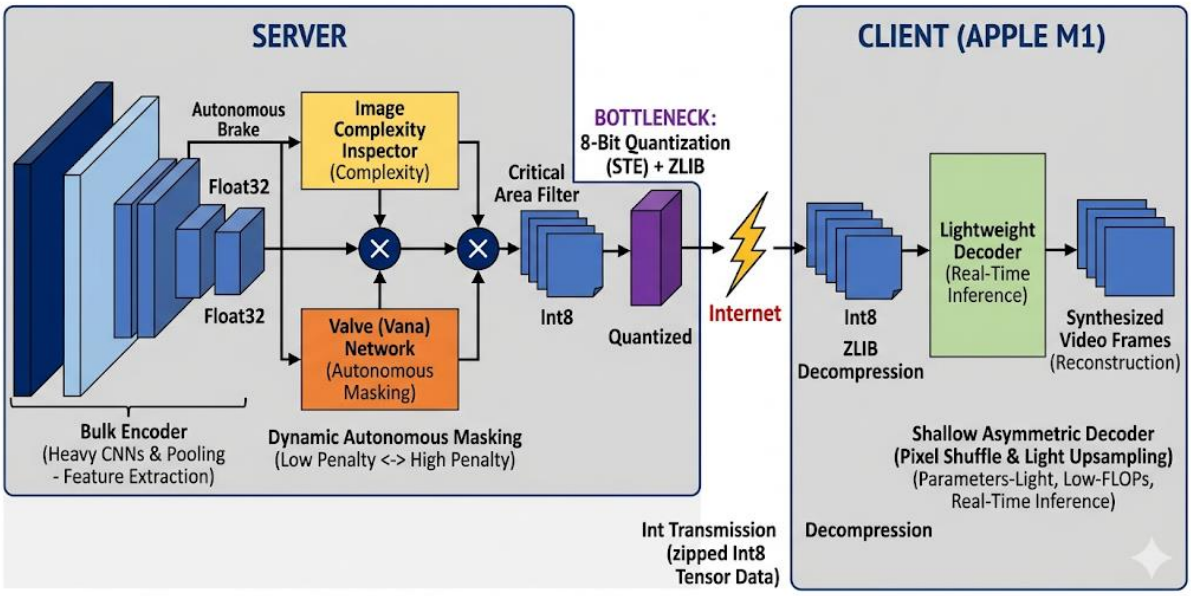

Figure 4.1: Technical block diagram of the end-to-end asymmetric split computing and autonomous masking architecture.

4.2. Bottleneck Theory, Quantization, and ZLIB Compression Compatibility

For the asymmetric architecture to provide bandwidth savings, it is mandatory to work integrated with classic dictionary-based lossless compression algorithms (ZLIB). However, in tests of the V3 model where spatial downsampling was not performed, it was discovered that the 16-channel floating-point (FP16) tensor data was structurally incompatible with the ZLIB algorithm. Due to continuously changing floating-point numbers, the original file size of 8.3 MB was measured as 54.3 MB (approximately 6.5 times) over the network.

To overcome this scientific crisis, the V5 Quantization Architecture was developed. Before being transmitted over the network, the tensor values at the bottleneck output were converted into an 8-bit discrete integer format using STE (Straight-Through Estimator) algorithms. Thanks to this manipulation, the data load of the ZLIB dictionary was lightened, and the transmission cost dropped to an average of 6.21 MB, falling below the original file size for the first time without spatial downsampling.

4.3. Autonomous Self-Checking Mechanism and the Valve Network

To further optimize transmission sizes, the "Valve" sub-network was designed. To maximize engineering efficiency, this valve mechanism was applied just before the image entered the heavy convolutional layers, enabling the filtering of critical areas. However, during the training process, it was detected that the artificial intelligence evaded the masking task by engaging in "Reward Hacking".

To solve this problem, the V6: Autonomous Self-Checking Model, which serves as an original contribution to the literature, was invented. Before compressing the data, the network measures the difficulty of the scene via an independent "Image Complexity Inspector". The resulting "Complexity Score" operates as a braking mechanism that the network applies to itself. The network forces masking by increasing the penalty coefficient in detail-less scenes such as the sky , while autonomously lowering the penalty in a complex forest scene, thereby ensuring the preservation of visual quality.

4.4. Experimental Findings and Visual Evidence

The success of the autonomous decision mechanism has been proven by the test statistics of the V6 model. In some scenes (e.g., 0834.png), it was observed that the zeroing (masking) rate rose to 37.0% levels, and the data, which had an original size of 8104 KB, dropped to 5454 KB. In the general average, even cases where the network autonomously eliminated 49.18% of the data resulted in a 0.943 SSIM performance.

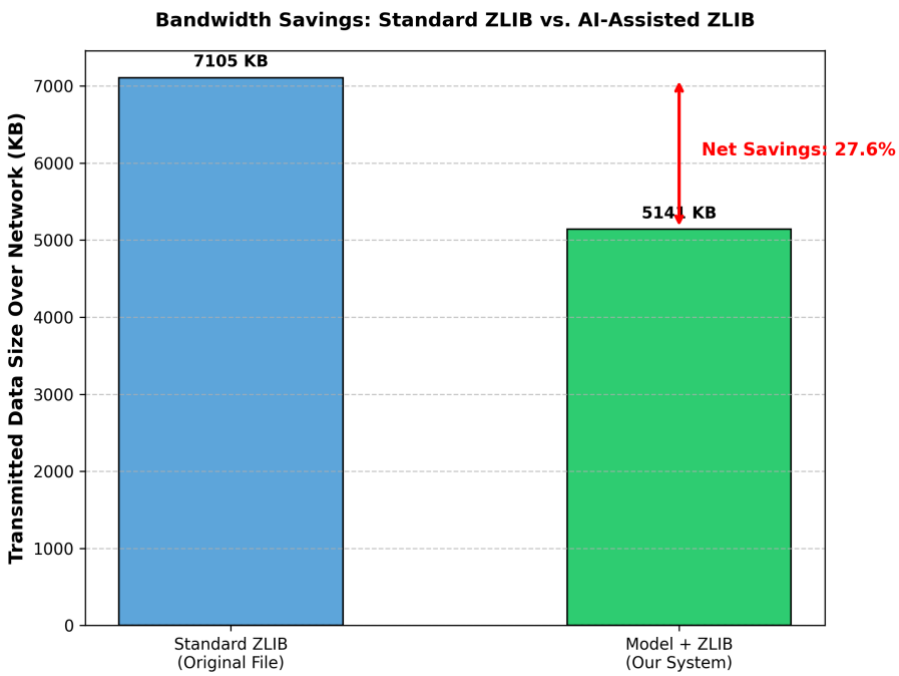

Figure 4.2: Comparison of standard ZLIB compression and model-assisted ZLIB compression on the basis of network traffic (KB).

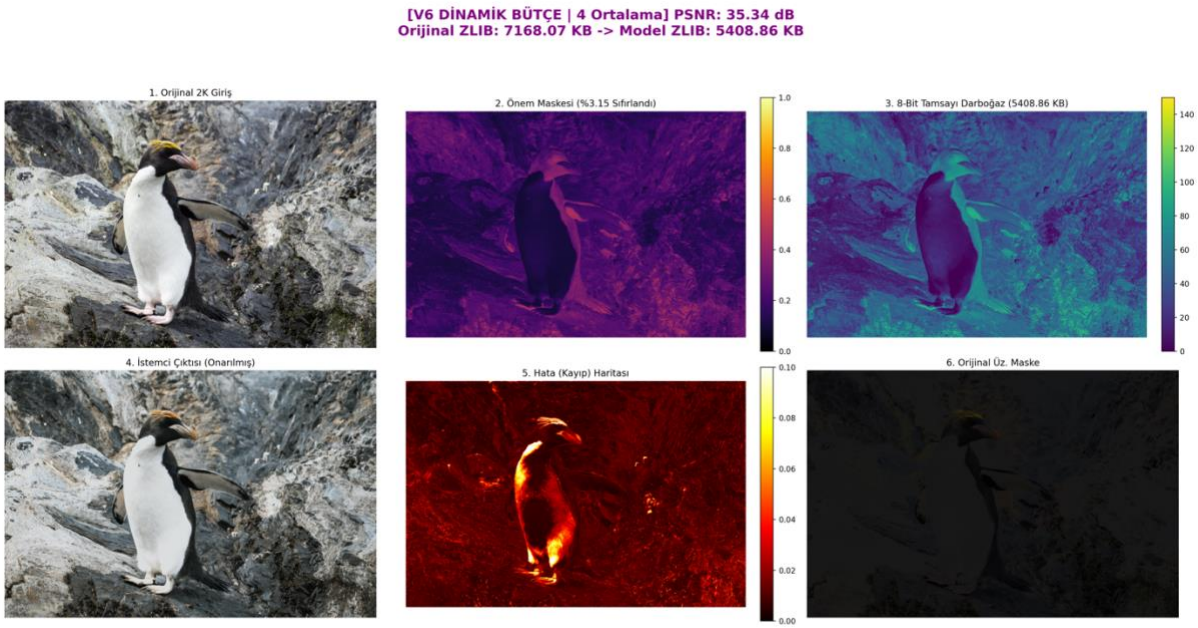

Figure 4.2: An average inference example where the V6 Autonomous Self-Checking architecture successfully applies its zeroing (masking) strategy and flawless visual integrity is maintained.

4.5. Color Fidelity and Chrominance Compression Challenges

Despite the high savings achieved with the V6 architecture, "color crushing" and losses in chrominance information were detected in the jury visuals (See Figure 4.3). The primary reason for this is that the image is processed as a single block in a 3-channel (RGB) space. The network, trying to pass through a narrow bottleneck, had to sacrifice color saturation to preserve the "skeleton" of the image and edge details.

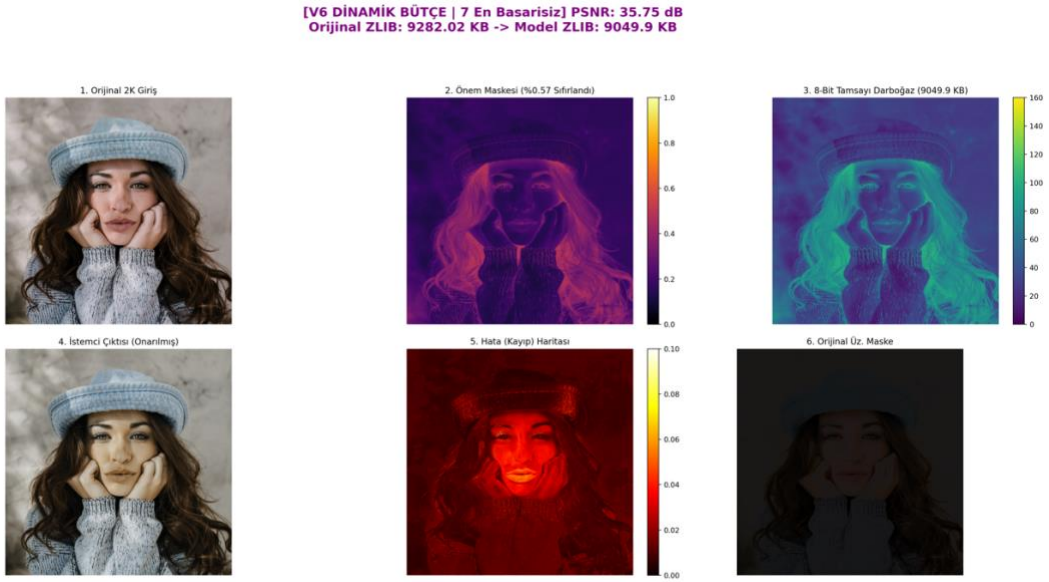

Figure 4.3: Color distortions resulting from the sacrifice of color information to structural details under aggressive masking (Unsuccessful example).

This "color crisis" became the fundamental source of motivation for the next phase of the project, Luma-Chroma separation (Separation of Concerns) , and cemented the vision that luminance and color information must be transmitted asymmetrically through different paths.

5. EXPLORING THE LIMITS: CRISES ENCOUNTERED IN ADVANCED BOTTLENECK OPTIMIZATION AND ENGINEERING LESSONS

The 38% autonomous bandwidth saving achieved with the V6 version of the project has been the most successful and stable point of the system in terms of practical application. However, Phase 5 (versions V8 - V18), initiated to push the model further, resolve color aberrations, and increase the compression ratio, did not yield the expected mathematical success. Nevertheless, this phase has been the most experience-gained, educational, and analytically deep R&D stage of the project in terms of understanding the nature of deep learning dynamics, control theory (PID), and data volume constraints.

In this section, the methods by which theoretical limits were pushed, the architectural crises encountered, and the engineering lessons learned from these failures are examined within an honest and objective framework.

5.1. Color Fading Crisis and the "Biological Isolation" Lesson (YUV)

The V6 model's sacrifice of colors to preserve structural texture (VGG) while performing aggressive masking created a bottleneck that could not be overcome as long as the work remained in the RGB space. In this crisis, where classical mathematical solutions (such as Cosine similarity) proved ineffective, one of the greatest engineering lessons was learned by drawing inspiration from nature.

Integrating the fact that the human eye is sensitive to brightness (Luma) but tolerant of color (Chroma), a transition was made from RGB to YUV space. While the Y (Luminance) channel was left to the VGG function, a strict L1 penalty was imposed on the U and V (Color) channels. Although this step increased the computational load on the network, it provided invaluable experience on how to construct the logic of "Perceptual Isolation" within artificial intelligence.

5.2. Capacity Paradox and Model Stubbornness (Reward Hacking)

One of the most significant obstacles encountered while attempting to compress the network further was the model's tendency to "evade its own penalty". When quantization shock, valve restriction, and quality penalties were applied to the network simultaneously, the model experienced "Vanishing Gradient," locked the valve at 100% open, and failed to function correctly.

This situation was diagnosed as the "Capacity Paradox," where a shallow network (45,159 parameters) ceases to operate effectively when it realizes it cannot reach demanding targets. As a solution, "Curriculum Learning" from the literature was implemented, and the training was divided into two distinct phases. Saturating the network with quality without limits first (the Painter phase), and then freezing the weights to command it strictly to compress the data (the Sculptor phase), represented a significant advancement for the project in understanding the fine-tuning of artificial intelligence training processes.

5.3. Dimensional Mismatch and Volume Conservation

Attempts at spatial narrowing (Downsampling) of the network resulted in tensor size mismatch errors between masks. This crisis was resolved by returning to fundamental mathematics rather than relying on standard library functions (stride=2). To ensure that 3-channel RGB data could maintain the same volume in a 6-channel bottleneck, the equation $a = b\sqrt{2}$ was established, and the network was made completely resolution-agnostic using F.interpolate. This solution was one of the most creative steps of the project in terms of understanding the physical equivalent of tensor mathematics.

5.4. The Ultimate Crisis: "Dimension Paradox" and Shaping the Future Vision

Despite all "Hard Rate Target" algorithms, the final point reached with the V18 model was a mathematical disappointment but a scientific enlightenment. Although the model successfully erased 59% of the pixels autonomously and maintained the quality at 30.29 dB PSNR, the ZLIB-compressed size of the masked data (averaging 14 MB) turned out to be double the original file size (averaging 7 MB).

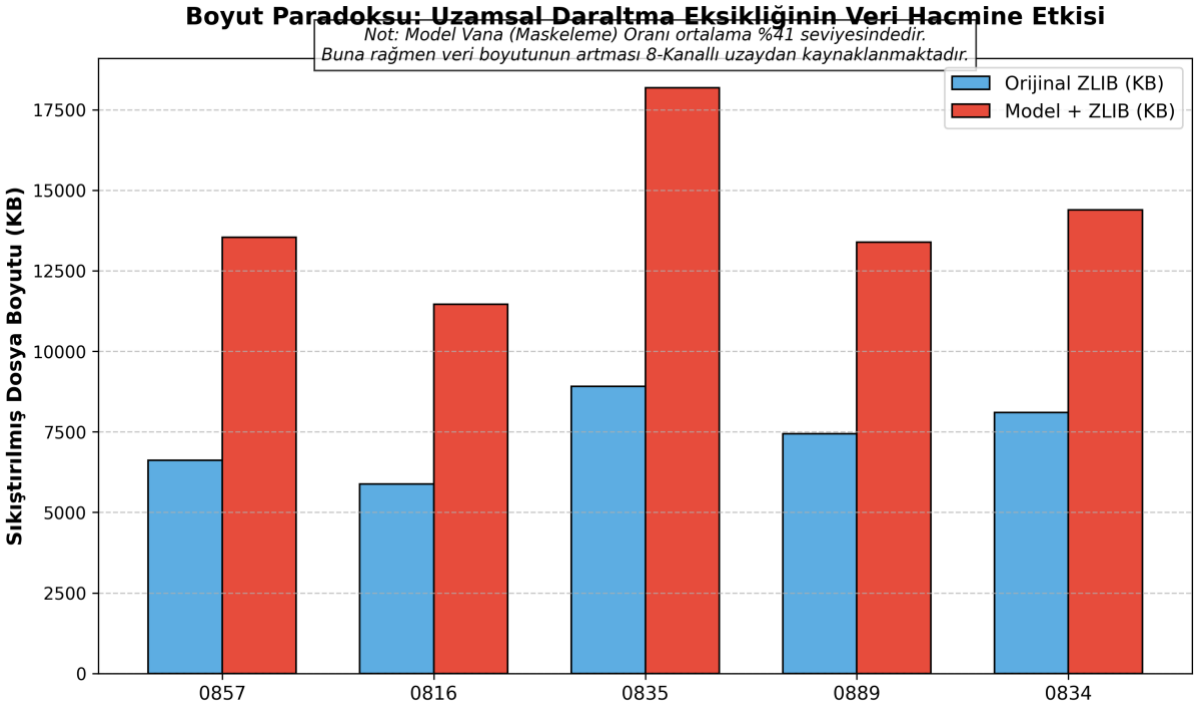

Figure 5.2: Physical data volume increase resulting from the lack of spatial narrowing, despite the model reducing the masking rate to 41%.

This result is not a "coding error" but a physical limit arising from the design philosophy of the architecture. To avoid increasing the computational cost (FLOPs), Spatial Downsampling was not utilized in the network. Transferring a 3-channel (RGB) matrix directly to an 8-channel latent space without reducing its aspect ratios expanded the data hardware-wise. Even the valve deleting half of the data could not compensate for this 8-channel physical bloating.

Section Conclusion and Inferences:

Although this phase fell behind V6 in terms of compression ratio, it represents the peak of the project regarding the technical insights gained. The "failure" mathematically proved that masking channels alone is insufficient in deep learning; physical narrowing of spatial dimensions is mandatory. This definitive diagnosis has clearly shaped the potential future "V19" vision of the project, which involves mandatory Downsampling on the Encoder side and Upsampling with PixelShuffle on the Decoder side.